

Feedback That Doesn't Backfire: Why a Third of Feedback Interventions Make Performance Worse

The best-replicated finding in feedback research is not that feedback works. It is that feedback is a double-edged sword — and the design of the conversation decides which edge lands.

The manager rehearsed the whole conversation on the drive in. The points were fair, the examples were real, the tone was kind. The meeting even ended with a handshake. Two weeks later the engineer's output was worse — slower, safer, more defensive — and the manager concluded the person "couldn't take feedback." The research suggests a different conclusion: the feedback was built to backfire, and it did exactly what feedback of that design usually does.

Ask almost any leadership book whether feedback improves performance and you get an unqualified yes. Ask the largest meta-analysis ever conducted on the question and you get something far more interesting — and far more useful.

Key takeaways

1. In Kluger and DeNisi's meta-analysis of 607 effect sizes covering 23,663 observations, feedback improved performance on average ($d = 0.41$), but more than one-third of feedback interventions — 38% of measured effects — *reduced* performance. [\[1\]](#)
2. Feedback intervention theory explains the reversal through attention: feedback effectiveness decreases as the feedback moves attention away from the task and toward the self. [\[1\]](#) [\[2\]](#)
3. Feedback works best when it is task-relevant, specific, compared against a clear goal or standard, and contains information about how to improve; cues that engage the self — including praise — attenuate its benefits. [\[1\]](#) [\[3\]](#)
4. In education research, Hattie and Timperley found feedback among the most powerful influences on achievement (average effect size 0.79) *and* among the most variable — the same double-edged pattern, in a different field. [\[3\]](#)
5. Feedback and goals are complementary mechanisms: goal-setting research shows feedback moderates goal effects, and feedback research shows goal-setting moderates feedback effects — neither works reliably without the other. [\[1\]](#) [\[5\]](#)
6. Employees are not passive receivers: feedback-seeking research shows people actively monitor and inquire for feedback as a resource, weighing its instrumental value against ego and image costs. [\[6\]](#) [\[7\]](#)
7. Applied to 360-degree appraisals, the same theory predicts — and the authors recommend designing for — task focus, goal linkage, and forward-looking development plans rather than person-focused verdicts. [\[4\]](#)

Does feedback improve performance?

On average, yes. Unconditionally, no — and the gap between those two answers is the most practically important finding in the literature.

In 1996, Avraham Kluger and Angelo DeNisi published a historical review and meta-analysis of feedback intervention research spanning nearly a century: 607 effect sizes, 23,663 observations. The average effect of feedback on performance was positive and moderate — $d = 0.41$. But the distribution around that average was the real news: **over one-third of feedback interventions (38% of measured effects) decreased performance.** [\[1\]](#)

Read that again the way a People leader should: if your managers deliver feedback the way the average intervention in a century of research delivered it, roughly one in three of those conversations is making performance worse. The variance is not noise. It is systematic, it is predictable, and Kluger and DeNisi built a theory to explain it.

This page is part of Cadence's [People Science](#) research pillar, and it exists because “give more feedback” is among the most common — and most under-specified — pieces of management advice in circulation. The evidence does not support giving more feedback. It supports giving *better-designed* feedback, and it is unusually clear about what the design is.

Why does feedback backfire?

Kluger and DeNisi's explanation — feedback intervention theory — is an attention model. [\[1\]](#) A feedback intervention changes the locus of a person's attention among three levels:

- **Task learning** — the details of how the work is done, and how to do it better.
- **Task motivation** — the gap between current performance and the goal or standard.
- **Meta-task, or self, processes** — what the feedback means about *me*: my competence, my standing, my identity.

The theory's central claim, elaborated in the authors' follow-up article, is that **feedback becomes less effective as attention climbs that hierarchy toward the self.** [\[2\]](#) Feedback that keeps attention on the task or on the task-standard gap tends to drive learning and effort. Feedback that engages the self diverts cognitive resources to managing emotion, defending self-image, or re-litigating the evaluation — resources that no longer go to the work. In the worst case the person disengages from the standard entirely: the goal is abandoned rather than pursued, because abandoning it protects the self.

Two details of this account are widely misread, and worth stating plainly:

Negative feedback is not the villain. The threat is *self-directed* feedback, in either direction. Discouraging feedback that reads as a verdict on the person backfires — and so, per the same mechanism, does person-directed praise. Both point attention at the self.

The backfire is not about feelings; it is about attention. Feedback intervention theory does not say uncomfortable feedback is harmful. It says feedback that spends the recipient's attention on the self rather than the task loses its mechanism of action. A hard, specific, task-focused message can be entirely safe under this theory. A warm, vague, person-focused one can be entirely counterproductive.

What separates feedback that works from feedback that hurts?

The meta-analytic moderators, plus the education literature that followed, converge on a short list. [\[1\]](#) [\[3\]](#)

Feedback tends to **help** when it:

- is about the **task or the process**, not the person;
- is **specific** — it identifies what happened and what better looks like, not a general impression;
- is anchored to a **clear goal or standard**, so the feedback describes a gap rather than a judgment;
- contains **information about how to improve** — a correct solution, a method, a next step;
- is **forward-looking** — it exists to shape the next attempt, not to adjudicate the last one.

Feedback tends to **hurt** when it:

- is **person-focused** — praise or criticism aimed at ability, character, or worth;
- is **vague** — “be more strategic,” “great job,” “needs more polish” — carrying no task information;
- **threatens the self** — public comparison, surprise verdicts, feedback experienced as evaluation of the person;
- is **normative rather than task-referenced** — ranking against peers instead of progress against a standard;
- arrives with **no path forward** — a gap with no method attached.

Hattie and Timperley’s synthesis of the education literature, “The Power of Feedback,” lands in the same place from independent data: feedback is among the most powerful influences on achievement — average effect size 0.79 across meta-analyses, roughly twice the typical education effect — and among the most variable. Their model distinguishes feedback about the task, the process, self-regulation, and the self, and finds feedback at the self level (praise of the person) the least effective of the four. Their organizing questions — *Where am I going? How am I going? Where to next?* — are a working definition of goal-anchored, forward-looking feedback. [\[3\]](#)

The goal-setting tradition supplies the complementary half. Locke and Latham’s summary of 35 years of goal research shows that specific, challenging goals outperform vague direction — provided feedback on progress is present. [\[5\]](#) Feedback needs a standard to compare against; goals need feedback to stay alive. The two literatures are, in effect, one mechanism described from opposite ends: a *task–standard gap made visible, with a route to close it*. See [Connecting Goals, Feedback, and Recognition](#) for how Cadence treats that loop.

Is praise safe?

Praise is the most counter-intuitive case, so it deserves its own answer.

Under feedback intervention theory, person-directed praise is a self-directed cue — it moves attention to the self exactly the way person-directed criticism does, and the meta-analysis found self-engaging cues attenuate feedback’s benefits. [\[1\]](#) Hattie and Timperley are blunter still: praise of the person is the least effective of their four feedback levels, because it carries no information about the task, the process, or what to do next. [\[3\]](#)

This is not an argument against recognizing people. It is an argument about *what the recognition is about*. Gallup’s workplace research emphasizes that effective recognition is honest, authentic, and individualized [\[8\]](#) — and specific recognition of a concrete contribution (“the way you sequenced the migration saved the launch”) is task-information wearing a positive sign. “You’re a star” is not. The practical rule that falls out of both literatures: **praise the work, specifically, not the person, generically.**

Does 360-degree feedback fix the problem?

Multi-rater feedback is often proposed as the corrective — more perspectives, less bias, harder to dismiss. DeNisi and Kluger examined 360-degree appraisals through the lens of feedback intervention theory and concluded that more feedback, from more directions, inherits the same double edge. [\[4\]](#) Multi-source feedback about *who you are*, delivered as a batch verdict, is close to a worst case for self-directed attention. Their recommendations are the familiar moderators applied to the 360 format: keep the content task-focused, tie it to goals and standards, and pair it with a forward-looking improvement plan rather than a score.

The delivery mechanism, in other words, matters as much as the data. Which raises the question of what the right delivery mechanism is.

Should managers wait for employees to ask?

One more finding rounds out the picture: employees are not passive targets of feedback. Ashford and Cummings showed four decades ago that people treat feedback as an individual resource and actively acquire it — by directly asking, and by monitoring the environment for cues. [\[6\]](#) The subsequent review literature maps the economics of that behavior: people weigh feedback's instrumental value against its ego cost (what I might hear) and image cost (what asking says about me). [\[7\]](#)

The design implication is underused: **if asking for feedback is costly, people stop asking, and the feedback that does arrive is the unsolicited, evaluative kind most likely to backfire.** A standing, private, recurring conversation drops the ego and image costs of feedback-seeking to near zero — asking “how am I doing on X?” inside a scheduled 1:1 is routine; asking it cold in a hallway is an event.

How does a structured 1:1 implement the moderators?

Here is Cadence's honest reading of its own product through this research. The claims below describe design intent mapped to evidence — not measured customer outcomes, because Cadence does not yet have customer outcome data and says so.

Cadence's structured 1:1 workspace — scheduling, agendas, shared notes, and action items — is **live today**, and it is deliberately shaped as a delivery mechanism for the moderators the research identifies:

- **Task anchor by default.** A structured agenda built from real work — goals, projects, prior action items — means feedback arrives attached to a task object, not as a free-floating verdict on the person. Goal and OKR tracking is live today, and goals surfaced in the conversation give feedback its standard: a gap to close, not a grade to absorb.
- **Expected, recurring, low-threat.** A weekly rhythm makes feedback a routine event inside an ongoing relationship rather than a surprise evaluation — the delivery condition least likely to trip self-directed attention, and the cheap feedback-seeking channel Ashford's literature implies employees need.
- **Forward-looking by construction.** Action items committed at the end of the conversation are the “where to next?” of Hattie and Timperley's model, made durable: the next 1:1 opens on what was agreed, so feedback becomes a loop with memory instead of an annual archaeology dig.
- **Specific recognition, not generic praise.** Cadence's recognition feed (live today) is built around recognizing concrete contributions — task-information with a positive sign — which is the form the research supports.
- **The manager stays the instrument.** Cadence's AI summaries and coaching prompts (preview, where enabled) exist to prepare a manager before the conversation and to remember commitments after it — nudging feedback toward task-specificity and follow-through. **AI develops managers, not replaces**

them; no AI in Cadence delivers feedback to an employee on a manager's behalf.

Managers, meanwhile, are the highest-leverage place to install this discipline: Gallup finds managers account for at least 70% of the variance in team engagement scores across business units [9], and Microsoft's Work Trend Index found 81% of employees say it is important for managers to help prioritize their workload while less than a third say managers have ever given clear one-on-one guidance. [10] The delivery mechanism is missing far more often than the willingness.

For the practitioner's playbook, see [How to Run Effective 1:1s](#); for the failure modes, [Why 1:1s Fail](#).

The strongest objection: does a 1996 meta-analysis of mostly lab tasks govern a 2026 workplace?

Steelman the skeptic, because the objection is serious:

Kluger and DeNisi's corpus leans heavily on laboratory and simple-task studies — memorization, vigilance, simple production tasks — many decades old. A manager's quarterly conversation about a senior engineer's architectural judgment is not a lab subject receiving trial-by-trial score feedback. The theory itself was labeled *preliminary* by its own authors. And "38% of effects were negative" is a statement about a heterogeneous research literature, not a base rate for your company's 1:1s. Isn't building product philosophy on this over-reading the evidence?

Three honest responses:

1. **The direction of the caveat favors more caution, not less.** Workplace feedback is *more* identity-laden than lab feedback — careers, compensation, and status ride on it — so the self-attention mechanism plausibly bites harder in organizations, not softer. The safe reading of the external validity gap is that the backfire risk is understated for the workplace, which argues for the moderators more strongly, not weakly.
2. **The pattern replicates across independent literatures.** Hattie and Timperley found the same signature — large average effects, huge variance, person-level feedback weakest — in education data with different tasks, populations, and decades. [3] The goal-setting literature independently confirms the feedback-standard coupling. [5] A mechanism that keeps reappearing across fields has earned more trust than any single corpus.
3. **The precise numbers are the weakest part of the claim; the design guidance is the strongest.** Cadence does not treat "38%" as a base rate for any real team, and neither should you. What the research licenses is a conditional: *unstructured, person-focused, standard-free feedback carries a materially non-trivial risk of harming performance, and a known set of design moderators reduces that risk*. That conditional — not the point estimate — is what Cadence's 1:1 design is built on.

What would change our mind: credible modern field evidence that person-focused and task-focused feedback perform equivalently in workplace settings would undercut the design thesis, and we would revise this page to say so.

What Cadence should not claim

Do not claim:

- "Cadence's 1:1s are scientifically proven to improve performance."
- "Feedback delivered through Cadence cannot backfire."

- “38% of your company’s feedback conversations are hurting performance.”
- “Cadence’s AI ensures feedback quality.”

Safer alternatives:

- “Feedback intervention research identifies design moderators; Cadence’s structured 1:1 workspace (live today) is built to implement them.”
- “Cadence makes task-anchored, goal-referenced, forward-looking feedback the default path, which is the form the evidence favors.”
- “Whether that design improves outcomes for Cadence customers is a future validation priority, not a current claim.”

How to cite this document

Suggested citation: Cadence, “Feedback That Doesn’t Backfire: Why a Third of Feedback Interventions Make Performance Worse” (2026). <https://cadencehr.ai/resources/feedback-that-doesnt-backfire>

Methodology and provenance. This synthesis was drafted in July 2026 by Cadence as part of the People Science research pillar. It is anchored on the Kluger–DeNisi feedback-intervention research program (the 1996 meta-analysis, the 1998 theoretical elaboration, and the 2000 application to multi-rater appraisal), triangulated against the independent education-research synthesis of Hattie and Timperley, goal-setting theory, and the feedback-seeking literature. Every citation below was verified to resolve to the named source as of 2026-07-27. Product-capability statements are labeled with current availability (live / preview / roadmap), and statements that are Cadence’s design interpretation rather than research findings are identified as such in the text.

References

1. Kluger, A. N., & DeNisi, A. (1996). “The Effects of Feedback Interventions on Performance: A Historical Review, a Meta-Analysis, and a Preliminary Feedback Intervention Theory.” *Psychological Bulletin*, 119(2), 254–284. [doi:10.1037/0033-2909.119.2.254](https://doi.org/10.1037/0033-2909.119.2.254)
2. Kluger, A. N., & DeNisi, A. (1998). “Feedback Interventions: Toward the Understanding of a Double-Edged Sword.” *Current Directions in Psychological Science*, 7(3), 67–72. [doi:10.1111/1467-8721.ep10772989](https://doi.org/10.1111/1467-8721.ep10772989)
3. Hattie, J., & Timperley, H. (2007). “The Power of Feedback.” *Review of Educational Research*, 77(1), 81–112. [doi:10.3102/003465430298487](https://doi.org/10.3102/003465430298487)
4. DeNisi, A. S., & Kluger, A. N. (2000). “Feedback Effectiveness: Can 360-Degree Appraisals Be Improved?” *Academy of Management Executive*, 14(1), 129–139. [doi:10.5465/ame.2000.2909845](https://doi.org/10.5465/ame.2000.2909845)
5. Locke, E. A., & Latham, G. P. (2002). “Building a Practically Useful Theory of Goal Setting and Task Motivation: A 35-Year Odyssey.” *American Psychologist*, 57(9), 705–717. [doi:10.1037/0003-066X.57.9.705](https://doi.org/10.1037/0003-066X.57.9.705)
6. Ashford, S. J., & Cummings, L. L. (1983). “Feedback as an Individual Resource: Personal Strategies of Creating Information.” *Organizational Behavior and Human Performance*, 32(3), 370–398. [doi:10.1016/0030-5073\(83\)90156-3](https://doi.org/10.1016/0030-5073(83)90156-3)
7. Ashford, S. J., Blatt, R., & VandeWalle, D. (2003). “Reflections on the Looking Glass: A Review of Research on Feedback-Seeking Behavior in Organizations.” *Journal of Management*, 29(6), 773–799. [doi:10.1016/S0149-2063\(03\)00079-5](https://doi.org/10.1016/S0149-2063(03)00079-5)
8. Gallup, “The Importance of Employee Recognition: Low Cost, High Impact.” gallup.com/workplace/236441

9. Gallup (2015). *State of the American Manager: Analytics and Advice for Leaders*; summarized in "Managers Account for 70% of Variance in Employee Engagement," Gallup Business Journal, April 2015. news.gallup.com/businessjournal/182792
10. Microsoft WorkLab, "Hybrid Work Is Just Work. Are We Doing It Wrong?" 2022 Work Trend Index Pulse Report. microsoft.com/worklab

This is a research synthesis, not a Cadence customer-outcome claim. Module availability is labeled because this page covers both live and preview capability.

This article is part of Cadence's [People Science](#) research pillar.

See how Cadence turns people science into operating rhythm at cadencehr.ai/product, or check plans at cadencehr.ai/pricing.